

## ▼ Import Libraries

The below block imports all needed libraries. Feel free to add additional libraries that you need and rerun below block.

Two last lines inform you of the Pytorch version and the availability of GPU. The last line should print GPU availability: True.

```
1 import os
2 import cv2
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import torch
6 import torch.nn as nn
7 import torch.nn.functional as F
8 import torch.optim as optim
9 from torch.utils.data import Dataset
10 import torchvision.transforms as transforms
11
12 print('Pytorch version: ', torch.__version__)
13 print('GPU availability: ', torch.cuda.is_available())
```

```
Pytorch version: 2.0.1+cu118
GPU availability: True
```

## ▼ Download Dataset

If you are familiar with Linux bash scripts, you can put `!` at the beginning of a command to order Colab of interpreting it as bash scripts instead of python scripts.

The below block downloads MNIST dataset and decompresses it.

```
1 !wget https://github.com/myleott/mnist_png/raw/master/mnist_png.tar.gz
2 !tar xzf mnist_png.tar.gz
```

```
--2023-10-03 20:16:51-- https://github.com/myleott/mnist_png/raw/master/mnist_png.ta
Resolving github.com (github.com)... 20.27.177.113
Connecting to github.com (github.com)|20.27.177.113|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://raw.githubusercontent.com/myleott/mnist_png/master/mnist_png.tar.gz
--2023-10-03 20:16:52-- https://raw.githubusercontent.com/myleott/mnist_png/master/r
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.110.133, 1
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.110.133|:
HTTP request sent, awaiting response... 200 OK
Length: 15683414 (15M) [application/octet-stream]
```

Saving to: 'mnist\_png.tar.gz'

mnist\_png.tar.gz 100%[=====>] 14.96M 41.5MB/s in 0.4s

2023-10-03 20:16:54 (41.5 MB/s) - 'mnist\_png.tar.gz' saved [15683414/15683414]

## Define Dataset Class

In order to use a dataset, we need to design a pytorch Dataset Class to process it. In the block below, you are required to complete:

- TODO1: def `__init__(root, transform)` function to build the MNIST dataset from images included in the `root` directory. Please add code below TODO1 to complete this function. The dataset should be captured by two lists, i.e., `self.images` that contains all images of MNIST, and `self.labels` that contains the corresponding label of each image in `self.images`.
- TODO2: def `__getitem__(index)` to draw a sample at `index` and its corresponding label. This function should return a tuple `(X, y)`, where `X` is the image (numpy ndarray of shape `1x28x28`) and `y` is a scalar from 0 to 9 representing `X`'s label.

```

1 # Define Dataset:
2 class MNISTDataset(Dataset):
3     def __init__(self, root, transform=None):
4         # `root` is expected to contain 10 sub-directories, each of which is named after
5         # transform is a torchvision.transforms object that pre-processes an image
6         self.root = root
7         self.transform = transform
8         #TODO 1: Read dataset.
9         # All images should be contained in a list `self.images`, and their correspondir
10        # `self.images[i]` should contain a numpy ndarray of size 1x28x28.
11        # `self.labels[i]` should contain a single integer of [0-9] representing the label
12        labels = os.listdir(root)
13        self.labels = []
14        self.images = []
15        for l in labels:
16            img_files = os.listdir(os.path.join(root, l))
17            for img_file in img_files:
18                img_file = os.path.join(root, l, img_file)
19                img = plt.imread(img_file)
20                self.images.append(img)
21                self.labels.append(int(l))
22
23    def __len__(self):
24        return len(self.labels)
25
```

```

26 def __getitem__(self, index):
27     # TODO2: retrieve `self.images[index]` and feed the image into self.transform.
28     # Then, return a tuple (X, y), where X is the image and y is its label.
29     image = self.images[index]
30     #image = self.transform(image)#transfor = self.transform(image)
31     label = self.labels[index]
32
33     if self.transform is not None:
34         image = self.transform(image)
35     return image, label
36
37 def show_random(self):
38     indices = np.random.randint(0, len(self), [16,])
39     f, ax = plt.subplots(4, 4, figsize=(10, 10))
40     for i in range(4):
41         for j in range(4):
42             ax[i, j].imshow(self.images[indices[i * 4 + j]])
43             ax[i, j].tick_params(top=False, bottom=False, left=False, right=False, :
44             ax[i, j].set_title(f'Label: {self.labels[indices[i * 4 + j]]}')
45     plt.show()

```

## Define Model Class

Below, we define a simple Multi-layer Perceptron Network with a hidden layer. A pytorch model necessarily have two functions, i.e., `__init__`, which defines all layers of the network, and `forward`, which is fed the input data and processes through all layers defined in `__init__`.

```

1 # Define Network:
2 class MLPNet(nn.Module):
3     def __init__(self):
4         super(MLPNet, self).__init__()
5         self.fc1 = nn.Linear(28*28, 500)
6         self.fc2 = nn.Linear(500, 256)
7         self.fc3 = nn.Linear(256, 10)
8     def forward(self, x):
9         x = x.view(-1, 28*28) # Flatten every image into a single vector
10        x = F.relu(self.fc1(x))
11        x = F.relu(self.fc2(x))
12        x = self.fc3(x)
13        return x
14
15    def name(self):
16        return "MLP"

```

## Create MNISTDataset objects and dataloaders

Below, we create the objects to process training and testing sets of MNIST data. As there are no held-out validation set, we manually split the training set into training and validation subsets with the ratio of 8:2.

After creating dataset objects, we wrap them by a Pytorch Dataloader to allow several necessary features in training deep learning models, e.g., mini-batch feeding, shuffling.

**Note:** if you successfully complete `__init__` function of `MNISTDataset`, its `show_random` function would successfully randomly show 16 images and corresponding labels in the dataset.

```

1 #####
2 # Hyper parameters
3 #####
4 BATCH_SIZE = 128
5 transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,), (1.0,))])
6
7 #####
8 # Create training and testing dataset and show random examples
9 #####
10 trainval_set = MNISTDataset('mnist_png/training', transform=transform)
11 trainval_set.show_random()
12
13 test_set = MNISTDataset('mnist_png/testing', transform=transform)
14
15 #####
16 # As there is no validation set
17 # We split training dataset into training and validation sets
18 #####
19 train_size = int(0.8 * len(trainval_set))
20 val_size = len(trainval_set) - train_size
21 train_set, val_set = torch.utils.data.random_split(
22     dataset=trainval_set,
23     lengths=[train_size, val_size],
24     generator=torch.Generator().manual_seed(42))
25
26 #####
27 # Print lengths of subsets
28 #####
29 print('Training set size: ', len(train_set))
30 print('Validation set size: ', len(val_set))
31 print('Testing set size: ', len(test_set))
32
33 #####
34 # Print lengths of subsets
35 #####
36 train_loader = torch.utils.data.DataLoader(
37     dataset=train_set,
38     batch_size=BATCH_SIZE,

```

```
38     batch_size=BATCH_SIZE,  
39     shuffle=True)  
40 val_loader = torch.utils.data.DataLoader(  
41     dataset=val_set,  
42     batch_size=BATCH_SIZE,  
43     shuffle=False)  
44 test_loader = torch.utils.data.DataLoader(  
45     dataset=test_set,  
46     batch_size=BATCH_SIZE,  
47     shuffle=False)
```

## Create Model and Training Process

In the below block, we create an object (our model) from the MLPNet deep neural network defined above.

Then, we create the `criterion` that will compute a loss value from predictions generated by our model and groundtruth labels. We also create the `optimizer`, which updates our model's learnable parameters based on the loss value to improve its performance.

Finally, we start training the model through `EPOCHS` number of epochs. At each epoch, after training the model through the training subset, we evaluate its loss and accuracy on validation subset. Usually, we would base on the loss or accuracy on validation subset to pick out the best performed model during our training process.

Your tasks:

- TODO 3: Based on average accuracy on validation set, save the model weights into a file. Hint: use `torch.save(model.state_dict(), PATH)` to save model weights into a file specified by `PATH`.
- TODO 4: Load the best model weights saved in `PATH` from above task into our `model`. Then, compute loss and accuracy of the best model on testing subset. Hint: use `checkpoint = torch.load(PATH)` to load content of file specified in `PATH` into `checkpoint`, then, use `model.load_state_dict(checkpoint)` to load parameters saved in `checkpoint` into `model`.

```

1 #####
2 # Hyper parameters
3 #####
4 LR = 0.001 # learning rate
5 EPOCHS = 100 # number of epochs to train model
6
7 #####
8 # Create model
9 #####
10 model = MLPNet().cuda()
11
12 #####
13 # Create optimizer and criterion
14 #####

```

```

15 optimizer = optim.SGD(model.parameters(), lr=LR, momentum=0.9)
16 criterion = nn.CrossEntropyLoss()
17
18 #####
19 # Training process
20 #####
21 for epoch in range(EPOCHS):
22     # training
23     total_loss = 0
24     for batch_idx, (x, target) in enumerate(train_loader):
25         optimizer.zero_grad()
26         x, target = x.cuda(), target.cuda()
27         out = model(x)
28         loss = criterion(out, target)
29         total_loss += loss.item()
30         loss.backward()
31         optimizer.step()
32     avg_loss = total_loss / len(train_set)
33     print(f'==>>> epoch: {epoch}, train loss: {avg_loss:.6f}')
34
35     # evaluating
36     correct_cnt, total_loss = 0, 0
37     for batch_idx, (x, target) in enumerate(val_loader):
38         x, target = x.cuda(), target.cuda()
39         out = model(x)
40         _, pred_label = torch.max(out, 1)
41         correct_cnt += (pred_label == target).sum()
42         # smooth average
43         total_loss += loss.item()
44     avg_loss = total_loss / len(val_set)
45     avg_acc = correct_cnt / len(val_set)
46     print(f'==>>> epoch: {epoch}, val loss: {avg_loss:.6f}, val accuracy: {avg_acc:.6f}')
47     # TODO3: Based on average accuracy on validation set, save the model weights into a
48     torch.save(model.state_dict(), '/content/model.pt')
49 #####
50 # Testing process
51 #####
52 # TODO4: use best performed model from the above process to compute loss and accuracy on
53 checkpoint = torch.load('/content/model.pt')
54 model.load_state_dict(checkpoint)

```

```

==>>> epoch: 0, train loss: 0.016876
==>>> epoch: 0, val loss: 0.014884, val accuracy: 0.629083
==>>> epoch: 1, train loss: 0.010452
==>>> epoch: 1, val loss: 0.006955, val accuracy: 0.817917
==>>> epoch: 2, train loss: 0.005341
==>>> epoch: 2, val loss: 0.004906, val accuracy: 0.860000
==>>> epoch: 3, train loss: 0.003926
==>>> epoch: 3, val loss: 0.003096, val accuracy: 0.876833
==>>> epoch: 4, train loss: 0.003385
==>>> epoch: 4, val loss: 0.004172, val accuracy: 0.887083
-->>> epoch: 5, train loss: 0.002087

```

```
==>>> epoch: 5, train loss: 0.003007
==>>> epoch: 5, val loss: 0.002875, val accuracy: 0.894500
==>>> epoch: 6, train loss: 0.002890
==>>> epoch: 6, val loss: 0.003184, val accuracy: 0.901667
==>>> epoch: 7, train loss: 0.002752
==>>> epoch: 7, val loss: 0.002455, val accuracy: 0.902917
==>>> epoch: 8, train loss: 0.002638
==>>> epoch: 8, val loss: 0.002466, val accuracy: 0.908750
==>>> epoch: 9, train loss: 0.002549
==>>> epoch: 9, val loss: 0.002003, val accuracy: 0.910583
==>>> epoch: 10, train loss: 0.002467
==>>> epoch: 10, val loss: 0.002012, val accuracy: 0.913333
==>>> epoch: 11, train loss: 0.002399
==>>> epoch: 11, val loss: 0.002777, val accuracy: 0.915167
==>>> epoch: 12, train loss: 0.002337
==>>> epoch: 12, val loss: 0.002475, val accuracy: 0.916583
==>>> epoch: 13, train loss: 0.002274
==>>> epoch: 13, val loss: 0.002487, val accuracy: 0.918083
==>>> epoch: 14, train loss: 0.002221
==>>> epoch: 14, val loss: 0.002206, val accuracy: 0.919417
==>>> epoch: 15, train loss: 0.002169
==>>> epoch: 15, val loss: 0.001963, val accuracy: 0.922167
==>>> epoch: 16, train loss: 0.002118
==>>> epoch: 16, val loss: 0.001949, val accuracy: 0.924750
==>>> epoch: 17, train loss: 0.002072
==>>> epoch: 17, val loss: 0.001545, val accuracy: 0.924667
==>>> epoch: 18, train loss: 0.002026
==>>> epoch: 18, val loss: 0.002137, val accuracy: 0.925833
==>>> epoch: 19, train loss: 0.001982
==>>> epoch: 19, val loss: 0.001700, val accuracy: 0.927917
==>>> epoch: 20, train loss: 0.001938
==>>> epoch: 20, val loss: 0.001926, val accuracy: 0.929917
==>>> epoch: 21, train loss: 0.001896
==>>> epoch: 21, val loss: 0.001522, val accuracy: 0.931167
==>>> epoch: 22, train loss: 0.001848
==>>> epoch: 22, val loss: 0.000932, val accuracy: 0.932500
==>>> epoch: 23, train loss: 0.001808
==>>> epoch: 23, val loss: 0.001848, val accuracy: 0.934000
==>>> epoch: 24, train loss: 0.001768
==>>> epoch: 24, val loss: 0.002049, val accuracy: 0.934667
==>>> epoch: 25, train loss: 0.001724
==>>> epoch: 25, val loss: 0.001141, val accuracy: 0.936583
==>>> epoch: 26, train loss: 0.001687
==>>> epoch: 26, val loss: 0.001747, val accuracy: 0.937167
==>>> epoch: 27, train loss: 0.001648
==>>> epoch: 27, val loss: 0.002331, val accuracy: 0.939417
==>>> epoch: 28, train loss: 0.001609
==>>> epoch: 28, val loss: 0.001549, val accuracy: 0.941417
```

## Training Famous State-of-the-art Neural Network on MNIST

In the next three blocks, you are requested to find the pytorch implementations for three famous state-of-the-art networks (i.e., LeNet, VGG16, and ResNet18) and train them using the training



process similar to the above block.

These tasks will help you have a comparison between state-of-the-arts. Your specific tasks are:

- TODO 5: Define LeNet network and train them using Cross Entropy loss, SGD optimizer, learning rate of 0.001, and in 100 epochs. Saving best performed model on validation subset during training process, and finally evaluate its performance (loss, accuracy) on testing set.
- TODO 6: Define VGG16 network and train them using Cross Entropy loss, SGD optimizer, learning rate of 0.001, and in 100 epochs. Saving best performed model on validation subset during training process, and finally evaluate its performance (loss, accuracy) on testing set.
- TODO 7: Define ResNet18 network and train them using Cross Entropy loss, SGD optimizer, learning rate of 0.001, and in 100 epochs. Saving best performed model on validation subset during training process, and finally evaluate its performance (loss, accuracy) on testing set.

```

1 # TODO5: Define LeNet network and train it using above training and testing processes
2 #transform.pad(2,2)
3
4 import torch
5 import torchvision
6 import numpy as np
7 import matplotlib.pyplot as plt
8
9 # From local helper files
10 #from helper_evaluation import set_all_seeds, set_deterministic, compute_confusion_matrix
11 #from helper_train import train_model
12 #from helper_plotting import plot_training_loss, plot_accuracy, show_examples, plot_confusion_matrix
13 #from helper_dataset import get_dataloaders_mnist
14
15
16 RANDOM_SEED = 123
17 BATCH_SIZE = 128
18 NUM_EPOCHS = 100
19 DEVICE = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
20
21 #set_all_seeds(RANDOM_SEED)

1 import torch.nn as nn
2 import torch.nn.functional as func
3
4 class lenet(nn.Module):
5     '''
6     input: 3x32x32 image

```

```

~      input: grayscale image
7      output: 10 class probability
8      '''
9      def __init__(self):
10         super(lenet, self).__init__()
11         self.conv1 = nn.Conv2d(3,6,5) #c1:featuremaps 6@28x28 #output = (input-filter)/:
12         self.conv2 = nn.Conv2d(6,16,5) #c3:feature_maps 16@10x10
13         self.maxPool = nn.MaxPool2d(2,2) #subsampling 1/2size
14         self.fc1 = nn.Linear(16*5*5,120) #f5:layer120
15         self.fc2 = nn.Linear(120,84) #f6:layer84
16         self.fc3 = nn.Linear(84,10) #output:10 class
17
18     def forward(self,x):
19         x = self.maxPool(func.relu(self.conv1(x)))
20         x = self.maxPool(func.relu(self.conv2(x)))
21         x = x.view(-1, 16*5*5) #flattens #tensor 형태 변환, 일렬로 16*5*5형태로 만들
22         x = func.relu(self.fc1(x))
23         x = func.relu(self.fc2(x))
24         x = self.fc3(x)
25         return x

1 #####
2 # Hyper parameters
3 #####
4 BATCH_SIZE = 128
5 transform = transforms.Compose([transforms.ToTensor(),transforms.Lambda(lambda x: x.repe
6
7 #####
8 # Create training and testing dataset and show random examples
9 #####
10 trainval_set = MNISTDataset('mnist_png/training', transform=transform)
11 trainval_set.show_random()
12
13 test_set = MNISTDataset('mnist_png/testing', transform=transform)
14
15 #####
16 # As there is no validation set
17 # We split training dataset into training and validation sets
18 #####
19 train_size = int(0.8 * len(trainval_set))
20 val_size = len(trainval_set) - train_size
21 train_set, val_set = torch.utils.data.random_split(
22     dataset=trainval_set,
23     lengths=[train_size, val_size],
24     generator=torch.Generator().manual_seed(42))
25
26 #####
27 # Print lengths of subsets
28 #####
29 print('Training set size: ', len(train_set))
30 print('Validation set size: ', len(val_set))

```

```
30 print('validation set size: ', len(val_set))
31 print('Testing set size: ', len(test_set))
32
33 #####
34 # Print lengths of subsets
35 #####
36 train_loader = torch.utils.data.DataLoader(
37     dataset=train_set,
38     batch_size=BATCH_SIZE,
39     shuffle=True)
40 val_loader = torch.utils.data.DataLoader(
41     dataset=val_set,
42     batch_size=BATCH_SIZE,
43     shuffle=False)
44 test_loader = torch.utils.data.DataLoader(
45     dataset=test_set,
46     batch_size=BATCH_SIZE,
47     shuffle=False)
```

```
1 #####
2 # Hyper parameters
3 #####
4 LR = 0.001 # learning rate
5 EPOCHS = 100 # number of epochs to train model
6
7 #####
8 # Create model
9 #####
10 model = lenet().cuda()
11
12 #####
13 # Create optimizer and criterion
14 #####
15 optimizer = optim.SGD(model.parameters(), lr=LR, momentum=0.9)
16 criterion = nn.CrossEntropyLoss()
17
18 #####
19 # Training process
20 #####
21 for epoch in range(EPOCHS):
22     # training
23     total_loss = 0
24     for batch_idx, (x, target) in enumerate(train_loader):
25         optimizer.zero_grad()
26         x, target = x.cuda(), target.cuda()
27         out = model(x)
28         loss = criterion(out, target)
29         total_loss += loss.item()
30         loss.backward()
31         optimizer.step()
32     avg_loss = total_loss / len(train_set)
33     print(f'==>> epoch: {epoch}, train loss: {avg_loss:.6f}')
34
35     # evaluating
36     #####
```

```

36     correct_cnt, total_loss = 0, 0
37     for batch_idx, (x, target) in enumerate(val_loader):
38         x, target = x.cuda(), target.cuda()
39         out = model(x)
40         _, pred_label = torch.max(out, 1)
41         correct_cnt += (pred_label == target).sum()
42         # smooth average
43         total_loss += loss.item()
44     avg_loss = total_loss / len(val_set)
45     avg_acc = correct_cnt / len(val_set)
46     print(f'==>> epoch: {epoch}, val loss: {avg_loss:.6f}, val accuracy: {avg_acc:.6f}')
47     # TODO3: Based on average accuracy on validation set, save the model weights into a
48     torch.save(model.state_dict(), '/content/model.pt')
49 #####
50 # Testing process
51 #####
52 # TODO4: use best performed model from the above process to compute loss and accuracy or
53 checkpoint = torch.load('/content/model.pt')
54 model.load_state_dict(checkpoint)

```

```

==>> epoch: 0, train loss: 0.017966
==>> epoch: 0, val loss: 0.018020, val accuracy: 0.168000
==>> epoch: 1, train loss: 0.017888
==>> epoch: 1, val loss: 0.017925, val accuracy: 0.195750
==>> epoch: 2, train loss: 0.017420
==>> epoch: 2, val loss: 0.016351, val accuracy: 0.606000
==>> epoch: 3, train loss: 0.008248
==>> epoch: 3, val loss: 0.003648, val accuracy: 0.867083
==>> epoch: 4, train loss: 0.002875
==>> epoch: 4, val loss: 0.002788, val accuracy: 0.915833
==>> epoch: 5, train loss: 0.002032
==>> epoch: 5, val loss: 0.001401, val accuracy: 0.935500
==>> epoch: 6, train loss: 0.001572
==>> epoch: 6, val loss: 0.000988, val accuracy: 0.945167
==>> epoch: 7, train loss: 0.001287
==>> epoch: 7, val loss: 0.001175, val accuracy: 0.955083
==>> epoch: 8, train loss: 0.001096
==>> epoch: 8, val loss: 0.000714, val accuracy: 0.959917
==>> epoch: 9, train loss: 0.000955
==>> epoch: 9, val loss: 0.000902, val accuracy: 0.964000
==>> epoch: 10, train loss: 0.000857
==>> epoch: 10, val loss: 0.000714, val accuracy: 0.967833
==>> epoch: 11, train loss: 0.000781
==>> epoch: 11, val loss: 0.000546, val accuracy: 0.970667
==>> epoch: 12, train loss: 0.000723
==>> epoch: 12, val loss: 0.000834, val accuracy: 0.973083
==>> epoch: 13, train loss: 0.000665
==>> epoch: 13, val loss: 0.000524, val accuracy: 0.974417
==>> epoch: 14, train loss: 0.000631
==>> epoch: 14, val loss: 0.000508, val accuracy: 0.972250
==>> epoch: 15, train loss: 0.000593
==>> epoch: 15, val loss: 0.000805, val accuracy: 0.974083
==>> epoch: 16, train loss: 0.000554
==>> epoch: 16, val loss: 0.000490, val accuracy: 0.977583

```

```
==>>> epoch: 17, train loss: 0.000531
==>>> epoch: 17, val loss: 0.000914, val accuracy: 0.978667
==>>> epoch: 18, train loss: 0.000512
==>>> epoch: 18, val loss: 0.000601, val accuracy: 0.976250
==>>> epoch: 19, train loss: 0.000483
==>>> epoch: 19, val loss: 0.000270, val accuracy: 0.977333
==>>> epoch: 20, train loss: 0.000468
==>>> epoch: 20, val loss: 0.000224, val accuracy: 0.981833
==>>> epoch: 21, train loss: 0.000445
==>>> epoch: 21, val loss: 0.000262, val accuracy: 0.980833
==>>> epoch: 22, train loss: 0.000430
==>>> epoch: 22, val loss: 0.000464, val accuracy: 0.980583
==>>> epoch: 23, train loss: 0.000417
==>>> epoch: 23, val loss: 0.000298, val accuracy: 0.980667
==>>> epoch: 24, train loss: 0.000408
==>>> epoch: 24, val loss: 0.000149, val accuracy: 0.981667
==>>> epoch: 25, train loss: 0.000384
==>>> epoch: 25, val loss: 0.000429, val accuracy: 0.981333
==>>> epoch: 26, train loss: 0.000371
==>>> epoch: 26, val loss: 0.000511, val accuracy: 0.981750
==>>> epoch: 27, train loss: 0.000363
==>>> epoch: 27, val loss: 0.000663, val accuracy: 0.982167
==>>> epoch: 28, train loss: 0.000354
==>>> epoch: 28, val loss: 0.000198, val accuracy: 0.983917
```

```
1 # TODO6: Define VGG16 network and train it using above training and testing processes
2 import torch.nn as nn
3 import torch.nn.functional as func
4 import math
5
6 class vggnet(nn.Module):
7     '''
8     input: 3x32x32 image
9     output: 10 class probability
10    '''
11    def __init__(self, features):
12        super(vggnet, self).__init__()
13        self.features = features
14        self.classifier = nn.Sequential(
15            nn.Dropout(),
16            nn.Linear(512 * 1 * 1, 512),
17            nn.ReLU(True),
18            nn.Dropout(),
19            nn.Linear(512, 512),
20            nn.ReLU(True),
21            nn.Linear(512, 10)
22        )
23        # Initialize weights
24        for m in self.modules():
25            if isinstance(m, nn.Conv2d):
26                n = m.kernel_size[0] * m.kernel_size[1] * m.out_channels
27                m.weight.data.normal_(0, math.sqrt(2. / n))
```

```
28         m.bias.data.zero_()
29
30     def forward(self, x):
31         x = self.features(x)
32         # print(sum(x[0, :]).detach().numpy()[0, 0],
33         #       sum(x[1, :]).detach().numpy()[0, 0],
34         #       sum(x[2, :]).detach().numpy()[0, 0],
35         #       sum(x[3, :]).detach().numpy()[0, 0])
36         x = x.view(-1, 512*1*1) #x = x.view(-1, x.size(0))
37         x = self.classifier(x)
38         # print(sum(x[0, :]).detach().numpy(),
39         #       sum(x[1, :]).detach().numpy(),
40         #       sum(x[2, :]).detach().numpy(),
41         #       sum(x[3, :]).detach().numpy())
42         return x
43
44 def make_layers(cfg):
45     layers = []
46     in_channels = 3
47     for v in cfg:
48         if v == 'M':
49             layers += [nn.MaxPool2d(kernel_size=2, stride=2)]
50         else:
51             if (v==257) or (v ==513):
52                 conv2d = nn.Conv2d(in_channels, v-1, kernel_size=1)
53                 in_channels = v-1
54             else:
55                 conv2d = nn.Conv2d(in_channels, v, kernel_size=3, padding=1)
56                 in_channels = v
57             layers += [conv2d, nn.ReLU(inplace=True)]
58
59     return nn.Sequential(*layers)
60
61 cfgs = {
62     'A': [64, 'M', 128, 'M', 256, 256, 'M', 512, 512, 'M', 512, 512, 'M'],
63     'B': [64, 64, 'M', 128, 128, 'M', 256, 256, 'M', 512, 512, 'M', 512, 512, 'M'],
64     'C': [64, 64, 'M', 128, 128, 'M', 256, 256, 257, 'M', 512, 512, 513, 'M', 512, 512,
65     'D': [64, 64, 'M', 128, 128, 'M', 256, 256, 256, 'M', 512, 512, 512, 'M', 512, 512,
66     'E': [64, 64, 'M', 128, 128, 'M', 256, 256, 256, 256, 'M', 512, 512, 512, 512, 'M',
67 }
68
69 def vgg11(): # configuration A
70     return vggnet(make_layers(cfgs['A']))
71
72 def vgg13(): # configuration B
73     return vggnet(make_layers(cfgs['B']))
74
75 def vgg16_1(): # configuration C
76     return vggnet(make_layers(cfgs['C']))
77
78 def vgg16(): # configuration D
```

```

79     return vggnet(make_layers(cfgs['D']))
80
81 def vgg19(): # configuration E
82     return vggnet(make_layers(cfgs['E']))

1 from collections import OrderedDict
2 class VGG(nn.Module):
3     """
4     Standard PyTorch implementation of VGG. Pretrained imagenet model is used.
5     """
6     def __init__(self):
7         super().__init__()
8
9         self.features = nn.Sequential(
10             # conv1
11             nn.Conv2d(3, 64, 3, padding=1),
12             nn.ReLU(),
13             nn.Conv2d(64, 64, 3, padding=1),
14             nn.ReLU(),
15             nn.MaxPool2d(2, stride=2, return_indices=True),
16
17             # conv2
18             nn.Conv2d(64, 128, 3, padding=1),
19             nn.ReLU(),
20             nn.Conv2d(128, 128, 3, padding=1),
21             nn.ReLU(),
22             nn.MaxPool2d(2, stride=2, return_indices=True),
23
24             # conv3
25             nn.Conv2d(128, 256, 3, padding=1),
26             nn.ReLU(),
27             nn.Conv2d(256, 256, 3, padding=1),
28             nn.ReLU(),
29             nn.Conv2d(256, 256, 3, padding=1),
30             nn.ReLU(),
31             nn.MaxPool2d(2, stride=2, return_indices=True),
32
33             # conv4
34             nn.Conv2d(256, 512, 3, padding=1),
35             nn.ReLU(),
36             nn.Conv2d(512, 512, 3, padding=1),
37             nn.ReLU(),
38             nn.Conv2d(512, 512, 3, padding=1),
39             nn.ReLU(),
40             nn.MaxPool2d(2, stride=2, return_indices=True),
41
42             # conv5
43             nn.Conv2d(512, 512, 3, padding=1),
44             nn.ReLU(),
45             nn.Conv2d(512, 512, 3, padding=1)

```



```

45         nn.Conv2d(512, 512, 3, padding=1),
46         nn.ReLU(),
47         nn.Conv2d(512, 512, 3, padding=1),
48         nn.ReLU(),
49         nn.MaxPool2d(2, stride=2, return_indices=True)
50     )
51
52     self.classifier = nn.Sequential(
53         nn.Linear(512 * 7 * 7, 4096),
54         nn.ReLU(),
55         nn.Dropout(),
56         nn.Linear(4096, 4096),
57         nn.ReLU(),
58         nn.Dropout(),
59         nn.Linear(4096, 1000)
60     )
61
62     # We need these for MaxUnpool operation
63     self.conv_layer_indices = [0, 2, 5, 7, 10, 12, 14, 17, 19, 21, 24, 26, 28]
64     self.feature_maps = OrderedDict()
65     self.pool_locs = OrderedDict()
66
67     def forward(self, x):
68         for layer in self.features:
69             if isinstance(layer, nn.MaxPool2d):
70                 x, location = layer(x)
71             else:
72                 x = layer(x)
73
74         x = x.view(x.size()[0], -1)
75         x = self.classifier(x)
76         return x
77
78
79 def get_vgg():
80     vgg = VGG()
81     temp = torchvision.models.vgg16(pretrained=True)
82     vgg.load_state_dict(temp.state_dict())
83     return vgg

```

  

```

1 #####
2 # Hyper parameters
3 #####
4 BATCH_SIZE = 128
5 transform = transforms.Compose([transforms.ToTensor(), transforms.Lambda(lambda x: x.repeat(
6
7 #####
8 # Create training and testing dataset and show random examples
9 #####
10 trainval_set = MNISTDataset('mnist_png/training', transform=transform)
11 trainval_loader = torch.utils.data.DataLoader(trainval_set, batch_size=BATCH_SIZE, shuffle=True)

```

```
11 trainval_set.snow_random()
12
13 test_set = MNISTDataset('mnist_png/testing', transform=transform)
14
15 #####
16 # As there is no validation set
17 # We split training dataset into training and validation sets
18 #####
19 train_size = int(0.8 * len(trainval_set))
20 val_size = len(trainval_set) - train_size
21 train_set, val_set = torch.utils.data.random_split(
22     dataset=trainval_set,
23     lengths=[train_size, val_size],
24     generator=torch.Generator().manual_seed(42))
25
26 #####
27 # Print lengths of subsets
28 #####
29 print('Training set size: ', len(train_set))
30 print('Validation set size: ', len(val_set))
31 print('Testing set size: ', len(test_set))
32
33 #####
34 # Print lengths of subsets
35 #####
36 train_loader = torch.utils.data.DataLoader(
37     dataset=train_set,
38     batch_size=BATCH_SIZE,
39     shuffle=True)
40 val_loader = torch.utils.data.DataLoader(
41     dataset=val_set,
42     batch_size=BATCH_SIZE,
43     shuffle=False)
44 test_loader = torch.utils.data.DataLoader(
45     dataset=test_set,
46     batch_size=BATCH_SIZE,
47     shuffle=False)
```

```
1 #####
2 # Hyper parameters
3 #####
4 LR = 0.001 # learning rate
5 EPOCHS = 100 # number of epochs to train model
6
7 #####
8 # Create model
9 #####
10 model = torchvision.models.vgg16().cuda()
11
12 #####
13 # Create optimizer and criterion
14 #####
15 optimizer = optim.SGD(model.parameters(), lr=LR, momentum=0.9)
16 criterion = nn.CrossEntropyLoss()
--
```

```
17
18 #####
19 # Training process
20 #####
21 for epoch in range(EPOCHS):
22     # training
23     total_loss = 0
24     for batch_idx, (x, target) in enumerate(train_loader):
25         optimizer.zero_grad()
26         x, target = x.cuda(), target.cuda()
27         out = model(x)
28         loss = criterion(out, target)
29         total_loss += loss.item()
30         loss.backward()
31         optimizer.step()
32     avg_loss = total_loss / len(train_set)
33     print(f'==>> epoch: {epoch}, train loss: {avg_loss:.6f}')
34
35     # evaluating
36     correct_cnt, total_loss = 0, 0
37     for batch_idx, (x, target) in enumerate(val_loader):
38         x, target = x.cuda(), target.cuda()
39         out = model(x)
40         _, pred_label = torch.max(out, 1)
41         correct_cnt += (pred_label == target).sum()
42         # smooth average
43         total_loss += loss.item()
44     avg_loss = total_loss / len(val_set)
45     avg_acc = correct_cnt / len(val_set)
46     print(f'==>> epoch: {epoch}, val loss: {avg_loss:.6f}, val accuracy: {avg_acc:.6f}')
47     # TODO3: Based on average accuracy on validation set, save the model weights into a
48     torch.save(model.state_dict(), '/content/model.pt')
49 #####
50 # Testing process
51 #####
52 # TODO4: use best performed model from the above process to compute loss and accuracy or
53 checkpoint = torch.load('/content/model.pt')
54 model.load_state_dict(checkpoint)
```

```
==>> epoch: 0, train loss: 0.022809
==>> epoch: 0, val loss: 0.018390, val accuracy: 0.112667
==>> epoch: 1, train loss: 0.016093
==>> epoch: 1, val loss: 0.007291, val accuracy: 0.619167
==>> epoch: 2, train loss: 0.003118
==>> epoch: 2, val loss: 0.001210, val accuracy: 0.949500
==>> epoch: 3, train loss: 0.001029
==>> epoch: 3, val loss: 0.001079, val accuracy: 0.964750
==>> epoch: 4, train loss: 0.000678
==>> epoch: 4, val loss: 0.000557, val accuracy: 0.976000
==>> epoch: 5, train loss: 0.000520
==>> epoch: 5, val loss: 0.000293, val accuracy: 0.977250
==>> epoch: 6, train loss: 0.000423
```

```

==>>> epoch: 6, val loss: 0.000228, val accuracy: 0.980917
==>>> epoch: 7, train loss: 0.000346
==>>> epoch: 7, val loss: 0.000327, val accuracy: 0.980167
==>>> epoch: 8, train loss: 0.000290
==>>> epoch: 8, val loss: 0.000331, val accuracy: 0.981083
==>>> epoch: 9, train loss: 0.000255
==>>> epoch: 9, val loss: 0.000035, val accuracy: 0.985083
==>>> epoch: 10, train loss: 0.000227
==>>> epoch: 10, val loss: 0.000844, val accuracy: 0.985250
==>>> epoch: 11, train loss: 0.000180
==>>> epoch: 11, val loss: 0.000026, val accuracy: 0.985000
==>>> epoch: 12, train loss: 0.000146
==>>> epoch: 12, val loss: 0.000113, val accuracy: 0.986833
==>>> epoch: 13, train loss: 0.000139
==>>> epoch: 13, val loss: 0.000099, val accuracy: 0.987500
==>>> epoch: 14, train loss: 0.000112
==>>> epoch: 14, val loss: 0.000451, val accuracy: 0.987333
==>>> epoch: 15, train loss: 0.000097
==>>> epoch: 15, val loss: 0.000144, val accuracy: 0.987750
==>>> epoch: 16, train loss: 0.000087
==>>> epoch: 16, val loss: 0.000164, val accuracy: 0.987333
==>>> epoch: 17, train loss: 0.000073
==>>> epoch: 17, val loss: 0.000022, val accuracy: 0.985500
==>>> epoch: 18, train loss: 0.000081
==>>> epoch: 18, val loss: 0.000386, val accuracy: 0.986333
==>>> epoch: 19, train loss: 0.000060
==>>> epoch: 19, val loss: 0.000009, val accuracy: 0.987500
==>>> epoch: 20, train loss: 0.000047
==>>> epoch: 20, val loss: 0.000086, val accuracy: 0.988167
==>>> epoch: 21, train loss: 0.000058
==>>> epoch: 21, val loss: 0.000171, val accuracy: 0.988250
==>>> epoch: 22, train loss: 0.000042
==>>> epoch: 22, val loss: 0.000139, val accuracy: 0.988250
==>>> epoch: 23, train loss: 0.000035
==>>> epoch: 23, val loss: 0.000001, val accuracy: 0.989167
==>>> epoch: 24, train loss: 0.000031
==>>> epoch: 24, val loss: 0.000004, val accuracy: 0.989417
==>>> epoch: 25, train loss: 0.000038
==>>> epoch: 25, val loss: 0.000001, val accuracy: 0.988250
==>>> epoch: 26, train loss: 0.000018
==>>> epoch: 26, val loss: 0.000012, val accuracy: 0.989917
==>>> epoch: 27, train loss: 0.000027
==>>> epoch: 27, val loss: 0.000021, val accuracy: 0.989083
==>>> epoch: 28, train loss: 0.000017
==>>> epoch: 28, val loss: 0.000002, val accuracy: 0.988750

```

```

1 # TODO7: Define Resnet18 network and train it using above training and testing processes
2 class ResBlock(nn.Module):
3     def __init__(self, in_channels, out_channels, downsample):
4         super().__init__()
5         if downsample:
6             self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3, stride=2, p
7             self.shortcut = nn.Sequential(
8                 nn.Conv2d(in_channels, out_channels, kernel_size=1, stride=2),

```

```
9         nn.BatchNorm2d(out_channels)
10     )
11     else:
12         self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3, stride=1, padding=1)
13         self.shortcut = nn.Sequential()
14
15     self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3, stride=1, padding=1)
16     self.bn1 = nn.BatchNorm2d(out_channels)
17     self.bn2 = nn.BatchNorm2d(out_channels)
18
19     def forward(self, input):
20         shortcut = self.shortcut(input)
21         input = nn.ReLU()(self.bn1(self.conv1(input)))
22         input = nn.ReLU()(self.bn2(self.conv2(input)))
23         input = input + shortcut
24         return nn.ReLU()(input)
25 class ResNet18(nn.Module):
26     def __init__(self, in_channels, resblock, outputs=1000):
27         super().__init__()
28         self.layer0 = nn.Sequential(
29             nn.Conv2d(in_channels, 64, kernel_size=7, stride=2, padding=3),
30             nn.MaxPool2d(kernel_size=3, stride=2, padding=1),
31             nn.BatchNorm2d(64),
32             nn.ReLU()
33         )
34
35         self.layer1 = nn.Sequential(
36             resblock(64, 64, downsample=False),
37             resblock(64, 64, downsample=False)
38         )
39
40         self.layer2 = nn.Sequential(
41             resblock(64, 128, downsample=True),
42             resblock(128, 128, downsample=False)
43         )
44
45         self.layer3 = nn.Sequential(
46             resblock(128, 256, downsample=True),
47             resblock(256, 256, downsample=False)
48         )
49
50
51         self.layer4 = nn.Sequential(
52             resblock(256, 512, downsample=True),
53             resblock(512, 512, downsample=False)
54         )
55
56         self.gap = torch.nn.AdaptiveAvgPool2d(1)
57         self.fc = torch.nn.Linear(512, outputs)
58
59     def forward(self, input):
```

```

55     def forward(self, input):
56         input = self.layer0(input)
57         input = self.layer1(input)
58         input = self.layer2(input)
59         input = self.layer3(input)
60         input = self.layer4(input)
61         input = self.gap(input)
62         input = torch.flatten(input)
63         input = self.fc(input)
64
65     return input

```

```

1 import torch
2 import torch.nn as nn
3
4
5 class Block(nn.Module):
6     def __init__(self, num_layers, in_channels, out_channels, identity_downsample=None,
7                 assert num_layers in [18, 34, 50, 101, 152], "should be a a valid architecture"
8                 super(Block, self).__init__()
9                 self.num_layers = num_layers
10                if self.num_layers > 34:
11                    self.expansion = 4
12                else:
13                    self.expansion = 1
14                # ResNet50, 101, and 152 include additional layer of 1x1 kernels
15                self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=1, stride=1, padding=0)
16                self.bn1 = nn.BatchNorm2d(out_channels)
17                if self.num_layers > 34:
18                    self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3, stride=1, padding=1)
19                else:
20                    # for ResNet18 and 34, connect input directly to (3x3) kernel (skip first (1x1) kernel)
21                    self.conv2 = nn.Conv2d(in_channels, out_channels, kernel_size=3, stride=1, padding=1)
22                self.bn2 = nn.BatchNorm2d(out_channels)
23                self.conv3 = nn.Conv2d(out_channels, out_channels * self.expansion, kernel_size=1, stride=1, padding=0)
24                self.bn3 = nn.BatchNorm2d(out_channels * self.expansion)
25                self.relu = nn.ReLU()
26                self.identity_downsample = identity_downsample
27
28    def forward(self, x):
29        identity = x
30        if self.num_layers > 34:
31            x = self.conv1(x)
32            x = self.bn1(x)
33            x = self.relu(x)
34            x = self.conv2(x)
35            x = self.bn2(x)
36            x = self.relu(x)
37            x = self.conv3(x)
38            x = self.bn3(x)
39
40        if self.identity_downsample is not None:
41            identity = self.identity_downsample(identity)
42
43        x = x + identity
44        x = self.relu(x)

```

```
39
40     if self.identity_downsample is not None:
41         identity = self.identity_downsample(identity)
42
43     x += identity
44     x = self.relu(x)
45     return x
46
47
48 class ResNet(nn.Module):
49     def __init__(self, num_layers, block, image_channels, num_classes):
50         assert num_layers in [18, 34, 50, 101, 152], f'ResNet{num_layers}: Unknown arch:
51                                                     f'to be 18, 34, 50, 101, or 152 '
52         super(ResNet, self).__init__()
53         if num_layers < 50:
54             self.expansion = 1
55         else:
56             self.expansion = 4
57         if num_layers == 18:
58             layers = [2, 2, 2, 2]
59         elif num_layers == 34 or num_layers == 50:
60             layers = [3, 4, 6, 3]
61         elif num_layers == 101:
62             layers = [3, 4, 23, 3]
63         else:
64             layers = [3, 8, 36, 3]
65         self.in_channels = 64
66         self.conv1 = nn.Conv2d(image_channels, 64, kernel_size=7, stride=2, padding=3)
67         self.bn1 = nn.BatchNorm2d(64)
68         self.relu = nn.ReLU()
69         self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)
70
71         # ResNetLayers
72         self.layer1 = self.make_layers(num_layers, block, layers[0], intermediate_channels)
73         self.layer2 = self.make_layers(num_layers, block, layers[1], intermediate_channels)
74         self.layer3 = self.make_layers(num_layers, block, layers[2], intermediate_channels)
75         self.layer4 = self.make_layers(num_layers, block, layers[3], intermediate_channels)
76
77         self.avgpool = nn.AdaptiveAvgPool2d((1, 1))
78         self.fc = nn.Linear(512 * self.expansion, num_classes)
79
80     def forward(self, x):
81         x = self.conv1(x)
82         x = self.bn1(x)
83         x = self.relu(x)
84         x = self.maxpool(x)
85
86         x = self.layer1(x)
87         x = self.layer2(x)
88         x = self.layer3(x)
89         x = self.layer4(x)
```



```

90
91     x = self.avgpool(x)
92     x = x.reshape(x.shape[0], -1)
93     x = self.fc(x)
94     return x
95
96     def make_layers(self, num_layers, block, num_residual_blocks, intermediate_channels):
97         layers = []
98
99         identity_downsample = nn.Sequential(nn.Conv2d(self.in_channels, intermediate_channels,
100                                                     nn.BatchNorm2d(intermediate_channels*self.expansion)
101         layers.append(block(num_layers, self.in_channels, intermediate_channels, identity_downsample))
102         self.in_channels = intermediate_channels * self.expansion # 256
103         for i in range(num_residual_blocks - 1):
104             layers.append(block(num_layers, self.in_channels, intermediate_channels)) #
105         return nn.Sequential(*layers)
106
107
108 def ResNet18(img_channels=3, num_classes=1000):
109     return ResNet(18, Block, img_channels, num_classes)
110
111
112 def ResNet34(img_channels=3, num_classes=1000):
113     return ResNet(34, Block, img_channels, num_classes)
114
115
116 def ResNet50(img_channels=3, num_classes=1000):
117     return ResNet(50, Block, img_channels, num_classes)
118
119
120 def ResNet101(img_channels=3, num_classes=1000):
121     return ResNet(101, Block, img_channels, num_classes)
122
123
124 def ResNet152(img_channels=3, num_classes=1000):
125     return ResNet(152, Block, img_channels, num_classes)
126
127
128 def test():
129     net = ResNet18(img_channels=3, num_classes=1000)
130     y = net(torch.randn(4, 3, 224, 224)).to("cuda")
131     print(y.size())
132
133
134 test()

```

```
torch.Size([4, 1000])
```

```

1 #####
2 # Hyper parameters

```

```
3 #####
4 BATCH_SIZE = 128
5 transform = transforms.Compose([transforms.ToTensor(), transforms.Lambda(lambda x: x.repeat(
6
7 #####
8 # Create training and testing dataset and show random examples
9 #####
10 trainval_set = MNISTDataset('mnist_png/training', transform=transform)
11 trainval_set.show_random()
12
13 test_set = MNISTDataset('mnist_png/testing', transform=transform)
14
15 #####
16 # As there is no validation set
17 # We split training dataset into training and validation sets
18 #####
19 train_size = int(0.8 * len(trainval_set))
20 val_size = len(trainval_set) - train_size
21 train_set, val_set = torch.utils.data.random_split(
22     dataset=trainval_set,
23     lengths=[train_size, val_size],
24     generator=torch.Generator().manual_seed(42))
25
26 #####
27 # Print lengths of subsets
28 #####
29 print('Training set size: ', len(train_set))
30 print('Validation set size: ', len(val_set))
31 print('Testing set size: ', len(test_set))
32
33 #####
34 # Print lengths of subsets
35 #####
36 train_loader = torch.utils.data.DataLoader(
37     dataset=train_set,
38     batch_size=BATCH_SIZE,
39     shuffle=True)
40 val_loader = torch.utils.data.DataLoader(
41     dataset=val_set,
42     batch_size=BATCH_SIZE,
43     shuffle=False)
44 test_loader = torch.utils.data.DataLoader(
45     dataset=test_set,
46     batch_size=BATCH_SIZE,
47     shuffle=False)
```

```
1 #####
2 # Hyper parameters
3 #####
4 LR = 0.001 # learning rate
5 EPOCHS = 100 # number of epochs to train model
6
7 #####
8 # Create model
```

```
9 #####
10 model = ResNet18().cuda()
11
12 #####
13 # Create optimizer and criterion
14 #####
15 optimizer = optim.SGD(model.parameters(), lr=LR, momentum=0.9)
16 criterion = nn.CrossEntropyLoss()
17
18 #####
19 # Training process
20 #####
21 for epoch in range(EPOCHS):
22     # training
23     total_loss = 0
24     for batch_idx, (x, target) in enumerate(train_loader):
25         optimizer.zero_grad()
26         x, target = x.cuda(), target.cuda()
27         out = model(x)
28         loss = criterion(out, target)
29         total_loss += loss.item()
30         loss.backward()
31         optimizer.step()
32     avg_loss = total_loss / len(train_set)
33     print(f'==>> epoch: {epoch}, train loss: {avg_loss:.6f}')
34
35     # evaluating
36     correct_cnt, total_loss = 0, 0
37     for batch_idx, (x, target) in enumerate(val_loader):
38         x, target = x.cuda(), target.cuda()
39         out = model(x)
40         _, pred_label = torch.max(out, 1)
41         correct_cnt += (pred_label == target).sum()
42         # smooth average
43         total_loss += loss.item()
44     avg_loss = total_loss / len(val_set)
45     avg_acc = correct_cnt / len(val_set)
46     print(f'==>> epoch: {epoch}, val loss: {avg_loss:.6f}, val accuracy: {avg_acc:.6f}')
47     # TODO3: Based on average accuracy on validation set, save the model weights into a
48     torch.save(model.state_dict(), '/content/model.pt')
49 #####
50 # Testing process
51 #####
52 # TODO4: use best performed model from the above process to compute loss and accuracy or
53 checkpoint = torch.load('/content/model.pt')
54 model.load_state_dict(checkpoint)
```

```
==>> epoch: 0, train loss: 0.002733
==>> epoch: 0, val loss: 0.000307, val accuracy: 0.981000
==>> epoch: 1, train loss: 0.000317
==>> epoch: 1, val loss: 0.000388, val accuracv: 0.983167
```

```
==>>> epoch: 2, train loss: 0.000164
==>>> epoch: 2, val loss: 0.000367, val accuracy: 0.984750
==>>> epoch: 3, train loss: 0.000077
==>>> epoch: 3, val loss: 0.000051, val accuracy: 0.985417
==>>> epoch: 4, train loss: 0.000046
==>>> epoch: 4, val loss: 0.000042, val accuracy: 0.987667
==>>> epoch: 5, train loss: 0.000027
==>>> epoch: 5, val loss: 0.000013, val accuracy: 0.987000
==>>> epoch: 6, train loss: 0.000018
==>>> epoch: 6, val loss: 0.000008, val accuracy: 0.987750
==>>> epoch: 7, train loss: 0.000012
==>>> epoch: 7, val loss: 0.000011, val accuracy: 0.988083
==>>> epoch: 8, train loss: 0.000010
==>>> epoch: 8, val loss: 0.000004, val accuracy: 0.987583
==>>> epoch: 9, train loss: 0.000007
==>>> epoch: 9, val loss: 0.000001, val accuracy: 0.988417
==>>> epoch: 10, train loss: 0.000007
==>>> epoch: 10, val loss: 0.000004, val accuracy: 0.988083
==>>> epoch: 11, train loss: 0.000005
==>>> epoch: 11, val loss: 0.000003, val accuracy: 0.988250
==>>> epoch: 12, train loss: 0.000004
==>>> epoch: 12, val loss: 0.000002, val accuracy: 0.988000
==>>> epoch: 13, train loss: 0.000004
==>>> epoch: 13, val loss: 0.000001, val accuracy: 0.988417
==>>> epoch: 14, train loss: 0.000004
==>>> epoch: 14, val loss: 0.000001, val accuracy: 0.987667
==>>> epoch: 15, train loss: 0.000004
==>>> epoch: 15, val loss: 0.000001, val accuracy: 0.987917
==>>> epoch: 16, train loss: 0.000003
==>>> epoch: 16, val loss: 0.000001, val accuracy: 0.988833
==>>> epoch: 17, train loss: 0.000004
==>>> epoch: 17, val loss: 0.000002, val accuracy: 0.988500
==>>> epoch: 18, train loss: 0.000003
==>>> epoch: 18, val loss: 0.000001, val accuracy: 0.988250
==>>> epoch: 19, train loss: 0.000003
==>>> epoch: 19, val loss: 0.000001, val accuracy: 0.988167
==>>> epoch: 20, train loss: 0.000003
==>>> epoch: 20, val loss: 0.000001, val accuracy: 0.988333
==>>> epoch: 21, train loss: 0.000003
==>>> epoch: 21, val loss: 0.000001, val accuracy: 0.988500
==>>> epoch: 22, train loss: 0.000002
==>>> epoch: 22, val loss: 0.000001, val accuracy: 0.988417
==>>> epoch: 23, train loss: 0.000002
==>>> epoch: 23, val loss: 0.000001, val accuracy: 0.987917
==>>> epoch: 24, train loss: 0.000002
==>>> epoch: 24, val loss: 0.000001, val accuracy: 0.988167
==>>> epoch: 25, train loss: 0.000002
==>>> epoch: 25, val loss: 0.000001, val accuracy: 0.988333
==>>> epoch: 26, train loss: 0.000002
==>>> epoch: 26, val loss: 0.000000, val accuracy: 0.988333
==>>> epoch: 27, train loss: 0.000002
==>>> epoch: 27, val loss: 0.000002, val accuracy: 0.988083
==>>> epoch: 28, train loss: 0.000002
==>>> epoch: 28, val loss: 0.000003, val accuracy: 0.988583
```

